

ASSIGNMENT-3
BIGDATA AND ECOSYSTEM
NAME: BHAVYA BHARDWAJ
BLEND ALL STARS

PySpark ETL Pipeline

Goal:

Align with the Python work to process large sales data using PySpark and generate enriched insights.

Tasks:

Read raw CSVs from HDFS or local folder.

Transform and clean the data (handle missing values, duplicates, etc.).

Enrich data by calculating KPIs such as:

Monthly Revenue

Profit Margin (%)

Region-wise Sales

Average Order Value

(Come up with 3–4 more KPIs as needed)

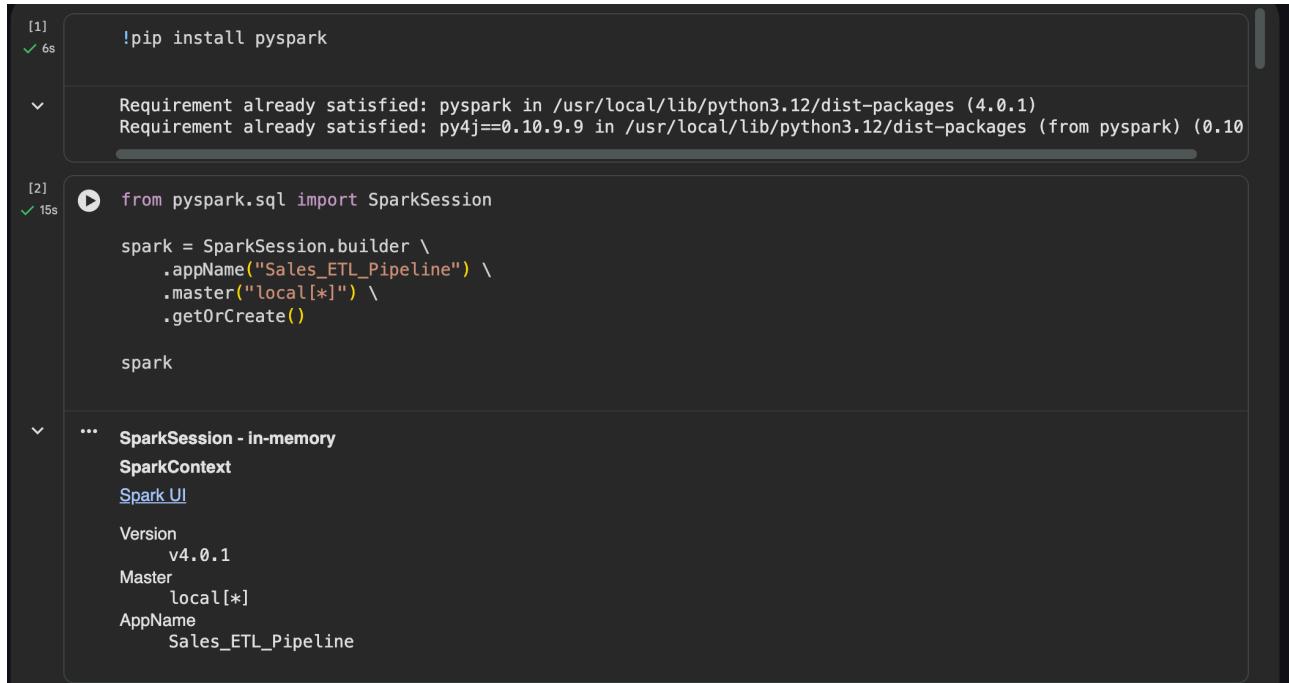
Write aggregated results to Parquet or a managed table.

(Optional) Integrate with Kafka for streaming order ingestion.

Tech: PySpark, HDFS, Kafka (optional)

Deliverables: PySpark script or notebook + pipeline diagram + screenshots (Spark job output, DAG view, Parquet output sample)

Dataset: https://wgcp-my.sharepoint.com/:f/g/personal/ritish_jogi_blend360_com/EmxzwFjNkaxPuCw2mQ0abr0BGg6XzlPlj22VogFVtQniyg?e=agISyr



The screenshot shows a Jupyter Notebook interface with two code cells. Cell [1] contains the command `!pip install pyspark`, which outputs that requirements are already satisfied. Cell [2] contains Python code to build a `SparkSession` with specific parameters like appName and master. The resulting `spark` object is then used to print its configuration details, showing it's an `in-memory` session with version v4.0.1, master set to `local[*]`, and appName set to `Sales_ETL_Pipeline`.

```
[1] 6s
!pip install pyspark

Requirement already satisfied: pyspark in /usr/local/lib/python3.12/dist-packages (4.0.1)
Requirement already satisfied: py4j==0.10.9.9 in /usr/local/lib/python3.12/dist-packages (from pyspark) (0.10

[2] 15s
▶ from pyspark.sql import SparkSession

spark = SparkSession.builder \
    .appName("Sales_ETL_Pipeline") \
    .master("local[*]") \
    .getOrCreate()

spark

...  SparkSession - in-memory
      SparkContext
      Spark UI

      Version           v4.0.1
      Master            local[*]
     AppName          Sales_ETL_Pipeline
```

```
m      from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

Ds    ⏎ import os

base_path = "/content/drive/MyDrive/blend_work_python/Sales Dataset"
print(os.listdir(base_path))

...  ['Amazon Sale Report.csv', 'Expense IIGF.csv', 'P L March 2021.csv', 'Sale Report.csv', 'Cloud Warehouse Com
```

```
[7]  ⏎ df = spark.read \
    .option("header", True) \
    .option("inferSchema", True) \
    .csv("/content/drive/MyDrive/blend_work_python/Sales Dataset/Amazon Sale Report.csv")

df.show(5)
df.printSchema()

... +---+-----+-----+-----+-----+-----+
|index|Order ID|Date|Status|Fulfilment|Sales Channel|ship-service-level|Style
+---+-----+-----+-----+-----+-----+
| 0|405-8078784-5731545|04-30-22|Cancelled|Merchant|Amazon.in|Standard|SET389
| 1|171-9198151-1101146|04-30-22|Shipped - Deliver...|Merchant|Amazon.in|Standard|JNE3781
| 2|404-0687676-7273146|04-30-22|Shipped|Amazon|Amazon.in|Expedited|JNE3371
| 3|403-9615377-8133951|04-30-22|Cancelled|Merchant|Amazon.in|Standard|J0341
| 4|407-1069790-7240320|04-30-22|Shipped|Amazon|Amazon.in|Expedited|JNE3671
+---+-----+-----+-----+-----+-----+
only showing top 5 rows
root
|-- index: integer (nullable = true)
|-- Order ID: string (nullable = true)
|-- Date: string (nullable = true)
|-- Status: string (nullable = true)
|-- Fulfilment: string (nullable = true)
|-- Sales Channel : string (nullable = true)
|-- ship-service-level: string (nullable = true)
|-- Style: string (nullable = true)
|-- SKU: string (nullable = true)
|-- Category: string (nullable = true)
|-- Size: string (nullable = true)
|-- ASIN: string (nullable = true)
|-- Courier Status: string (nullable = true)
|-- Qty: integer (nullable = true)
|-- currency: string (nullable = true)
```

Basic data cleaning

```
[9]  ⏎ df = df.drop("index", "Unnamed: 22")

[10] ⏎ df = df.withColumnRenamed("Sales Channel ", "Sales_Channel")

[11] ⏎ #basic data cleaning
     df = df.dropDuplicates()

[12] ⏎ df = df.dropna(subset=["Order ID", "Date", "Qty", "Amount"])

[13] ⏎ df = df.fillna({
    "ship-city": "Unknown",
    "ship-state": "Unknown",
    "promotion-ids": "None",
    "fulfilled-by": "Unknown"
})

[14] ⏎ from pyspark.sql.functions import col
     df = df.withColumn("Qty", col("Qty").cast("int")) \
             .withColumn("Amount", col("Amount").cast("double"))
```

BASIC KPI :DELIVERABLES NEEDED

```
[ ] #KPI CALCULATIONS

[19] ⏪ #monthly revenue
    7s

        from pyspark.sql.functions import year, month, sum

        monthly_revenue = df.withColumn("Year", year("Order_Date")) \
            .withColumn("Month", month("Order_Date")) \
            .groupBy("Year", "Month") \
            .agg(sum("Amount").alias("Monthly_Revenue"))

        monthly_revenue.show()

    ... +---+---+
        |Year|Month| Monthly_Revenue|
    +---+---+
        |2022| 3| 101683.85|
        |2022| 5| 2.622500474999955E7|
        |2022| 6| 2.342464637999988E7|
        |2022| 4| 2.883870832000023E7|
    +---+---+
```

```
] ⏪ #profit margin
    ss

        from pyspark.sql.functions import col, sum

        df = df.withColumn("Cost", col("Amount") * 0.7)
        df = df.withColumn("Profit", col("Amount") - col("Cost"))

        profit_margin_df = df.agg(
            (sum("Profit") / sum("Amount") * 100).alias("Profit_Margin_Percentage")
        )

        profit_margin_df.show()

    ... +---+
        |Profit_Margin_Percentage|
    +---+
        | 30.00000000000917|
    +---+
```

```
[21] #region wise sales
region_sales_df = df.groupBy("ship-state") \
    .agg(sum("Amount").alias("Total_Sales")) \
    .orderBy(col("Total_Sales").desc())
region_sales_df.show()

...
+-----+-----+
| ship-state| Total_Sales|
+-----+-----+
| MAHARASHTRA| 1.3334595140000008E7|
| KARNATAKA| 1.0481114370000001E7|
| TELANGANA| 6916615.650000002|
| UTTAR PRADESH| 6816109.08|
| TAMIL NADU| 6515650.110000002|
| DELHI| 4235215.97|
| KERALA| 3830227.58|
| WEST BENGAL| 3507880.440000004|
| ANDHRA PRADESH| 3219831.719999997|
| HARYANA| 2882092.989999998|
| Gujarat| 2728651.820000003|
| RAJASTHAN| 1716802.400000001|
| MADHYA PRADESH| 1592382.980000002|
| BIHAR| 1394388.319999998|
| ODIsha| 1372205.63|
| PUNJAB| 1180064.840000003|
| ASSAM| 1018136.2|
| UTTARAKHAND| 974143.55|
| JHARKHAND| 919088.21|
| GOA| 619437.85|
+-----+
only showing top 20 rows
```

```
[22] #avg order value
from pyspark.sql.functions import avg
aov_df = df.agg(
    avg("Amount").alias("Average_Order_Value")
)
aov_df.show()

...
+-----+
|Average_Order_Value|
+-----+
| 648.5557762611699|
+-----+
```

```
#extra KPI's MONTHLY ORDER COUNT
from pyspark.sql.functions import year, month, count

monthly_orders_df = df.groupBy(
    year("Order_Date").alias("Year"),
    month("Order_Date").alias("Month")
).agg(
    count("Order ID").alias("Total_Orders")
)

monthly_orders_df.show()

...
+---+---+---+
|Year|Month|Total_Orders|
+---+---+---+
|2022| 3| 162|
|2022| 5| 39534|
|2022| 6| 35413|
|2022| 4| 46068|
+---+---+---+
```

[24] #top selling categories by quantity

```
top_categories_df = df.groupBy("Category") \
    .agg(sum("Qty").alias("Total_Quantity")) \
    .orderBy(col("Total_Quantity").desc())

top_categories_df.show(5)
```

...

| Category | Total_Quantity |
|---------------|----------------|
| Set | 45223 |
| kurta | 44969 |
| Western Dress | 13939 |
| Top | 9899 |
| Ethnic Dress | 1053 |

only showing top 5 rows

STORING THE OPTIMIZED RESULTS IN PARQUET

[27] #the results optimised and stored in parquet for faster columunar retrieval of the KPIs ,preferred by bug cor

```
base_output_path = "/content/drive/MyDrive/blend_work_python/output/amazon_kpis"

profit_margin_df.write.mode("overwrite").parquet(f"{base_output_path}/profit_margin")
region_sales_df.write.mode("overwrite").parquet(f"{base_output_path}/region_sales")
aov_df.write.mode("overwrite").parquet(f"{base_output_path}/average_order_value")
monthly_orders_df.write.mode("overwrite").parquet(f"{base_output_path}/monthly_orders")
top_categories_df.write.mode("overwrite").parquet(f"{base_output_path}/top_categories")
```

[28] #to check if our parquet folder exists

```
import os

output_path = "/content/drive/MyDrive/blend_work_python/output/amazon_kpis/region_sales"
print(os.listdir(output_path))

['part-00000-27b68ffa-060b-4c62-ac36-2f0fba99234c-c000.snappy.parquet', '.part-00000-27b68ffa-060b-4c62-ac36-2f0fba99234c-c000.snappy.parquet')
```

[29] parquet_df = spark.read.parquet("/content/drive/MyDrive/blend_work_python/output/amazon_kpis/region_sales")
parquet_df.show(5)

...

| ship-state | Total_Sales |
|---------------|----------------------|
| MAHARASHTRA | 1.3334595140000008E7 |
| KARNATAKA | 1.0481114370000001E7 |
| TELANGANA | 6916615.650000002 |
| UTTAR PRADESH | 6816109.08 |
| TAMIL NADU | 6515650.110000002 |

only showing top 5 rows

DATA ETL PIPELINE DIAGRAM



PS: WE DID NOT USE HDFS AS OUR DATASET WAS SMALL AND DID NOT REQUIRE HDFS AND COULD EASILY RUN ON GOOGLE COLLAB USING LOCAL FILE STORAGE